

# Gesture Detection Models with VGG16 Base

Richeek Das

May 4, 2020

## 1 Introduction

I was playing around with **Transfer Learning** for my Hand Gesture Detection module, using pretrained convnets like *VGG16*, *ResNet50*, *GoogleNET*, *DenseNet121* while exploring and fine-tuning the hyperparameters to deliver dependable results in real-life applications. So here is a detailed specification of the models with a **VGG16** convolutional base that I had used, along with the *Accuracy and Loss Curves*, *Confusion Matrices*, *F1 scores and Error Rates*.

## 2 Dataset Description

I had 550 different images for each class spread over 5 classes :

Train Set : 375 images per class

Validation Set : 125 images per class

Test Set : 50 images per class (It's pretty less though)

### 2.1 Data Augmentation used :

```
rescale=1./255
rotation_range=45
width_shift_range=0.2
height_shift_range=0.2
shear_range=0.2
zoom_range=0.2
horizontal_flip=True
fill_mode='nearest'
```

A sample of the dataset(Just to check what it looks like) :



### 3 Models

**Convolution Base** has been taken to be the **VGG16** model with an input size of **(224,224,3)** with pretrained weights from **"imagenet"** training set.

#### 3.1 Model 1

Just to get a working baseline with our dataset.

##### 3.1.1 Specifications

**Model Classifier Top :**

```
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(5, activation='softmax'))
```

**Model Compilation :**

```
optimizer = Adam in default
loss = "categorical_crossentropy"
```

**Model Fitting :**

```
epochs = 50
steps_per_epoch = 100
batch_size = 15
validation_steps = 50
```

### 3.1.2 Results

*Validation Accuracy : 90%*

*F1 Score : 0.8864*


#### Confusion Matrix :


Class Indices : {'cool': 0, 'fist': 1, 'ok': 2, 'stop': 3, 'yo': 4}

```
array([[26,  0,  4,  3, 17],
       [ 0, 48,  0,  2,  0],
       [ 0,  0, 50,  0,  0],
       [ 0,  0,  0, 50,  0],
       [ 0,  0,  0,  0, 50]], dtype=int64)
```

Seems like it is pretty heavily confusing, “cool” sign with “yo” which is kind of apparently obvious, since both of them have 2 fingers

### 3.1.3 Accuracy and Loss Curves





## 3.2 Model 2

Decreased the *learning rate* of the optimizer, increased the *train batch size* from 15 to 32, set the *validation steps* and *steps per epoch* to something that seemed optimal and increased the *no of epochs* to analyse it further.

### 3.2.1 Specifications

#### Model Classifier Top :

```
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(5, activation='softmax'))
```

#### Model Compilation :

```
optimizer = Adam(lr=2e-5)
loss = "categorical_crossentropy"
```

#### Model Fitting :

```
epochs = 100
steps_per_epoch = 120
batch_size = 32
validation_steps = None
```

### 3.2.2 Results

*Validation Accuracy : 93%*

*F1 Score : 0.9207*


#### Confusion Matrix :


Class Indices : {'cool': 0, 'fist': 1, 'ok': 2, 'stop': 3, 'yo': 4}

```
array([[35,  0,  4,  4,  7],
       [ 0, 48,  0,  2,  0],
       [ 0,  0, 50,  0,  0],
       [ 0,  0,  0, 50,  0],
       [ 0,  0,  0,  2, 48]], dtype=int64)
```

Seems like it is still heavily confusing, “cool” sign with “yo”.

### 3.2.3 Accuracy and Loss Curves





### 3.3 Model 3

Increased the *learning rate* of the optimizer, did major changes to the model classifier(as listed below) and increased the *no of epochs* to analyse it further.

#### 3.3.1 Specifications

##### Model Classifier Top :

```
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(5, activation='softmax'))
```

##### Model Compilation :

```
optimizer = Adam(lr=1e-4)
loss = "categorical_crossentropy"
```

##### Model Fitting :

```
epochs = 200
steps_per_epoch = 120
batch_size = 32
validation_steps = None
```

#### 3.3.2 Results

Model is **UNDERFITTING** and the *loss* is still pretty high :(

*Validation Accuracy : 95%*

F1 Score : 0.9429

### Confusion Matrix :


Class Indices : {'cool': 0, 'fist': 1, 'ok': 2, 'stop': 3, 'yo': 4}


```
array([[40,  0,  3,  2,  5],
       [ 0, 49,  0,  1,  0],
       [ 0,  0, 50,  0,  0],
       [ 0,  0,  0, 50,  0],
       [ 0,  0,  0,  3, 47]], dtype=int64)
```

Seems like it is still confusing, “cool” sign with “yo” :(

But first, we need to fight the Underfit problem !

### 3.3.3 Accuracy and Loss Curves





### 3.4 Model 4

Removed the *Dropout layer* from the classifier top and further increased the *no of epochs* to analyse it further.

#### 3.4.1 Specifications

##### Model Classifier Top :

```
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(5, activation='softmax'))
```

##### Model Compilation :

```
optimizer = Adam(lr=1e-4)
loss = "categorical_crossentropy"
```

##### Model Fitting :

```
epochs = 200
steps_per_epoch = 120
batch_size = 32
validation_steps = None
```

#### 3.4.2 Results

*Validation Accuracy : 96.8%*

*F1 Score : 0.9344*



### Confusion Matrix :

Class Indices : {'cool': 0, 'fist': 1, 'ok': 2, 'stop': 3, 'yo': 4}


```
array([[38,  0,  2,  4,  6],
       [ 0, 47,  0,  3,  0],
       [ 0,  0, 50,  0,  0],
       [ 0,  0,  0, 50,  0],
       [ 0,  0,  0,  1, 49]], dtype=int64)
```


Seems like it is still confusing, “cool” sign with “yo” -\_-

But, as we see the Underfit problem has been solved !

Now let’s fine tune the final layers of the freezed **VGG16** base to gain a performance boost !

#### 3.4.3 Accuracy and Loss Curves





### 3.5 Model 4 Fine Tuned

Unfreezed block5\_conv1 layer of vgg16 and retrained it with a reduced learning rate :

```
epochs = 100
optimizer = optimizers.Adam(lr=1e-4)
loss = "categorical_crossentropy"
```

#### 3.5.1 Results

*Validation Accuracy : 98.08%*

*F1 Score = 0.9959*

#### Confusion Matrix :

Class Indices : {'cool': 0, 'fist': 1, 'ok': 2, 'stop': 3, 'yo': 4}

```
array([[50, 0, 0, 0, 0],
       [ 0, 50, 0, 0, 0],
       [ 0, 0, 50, 0, 0],
       [ 0, 1, 0, 49, 0],
       [ 0, 0, 0, 0, 50]], dtype=int64)
```

Seems like that fixed its confusion :)

### 3.5.2 Accuracy and Loss Curves

